

Exploration of DuckDB

JORRIT VANDER MYNSBRUGGE

GROUP 14

SID 0606134


WORKING STUDENT

Context

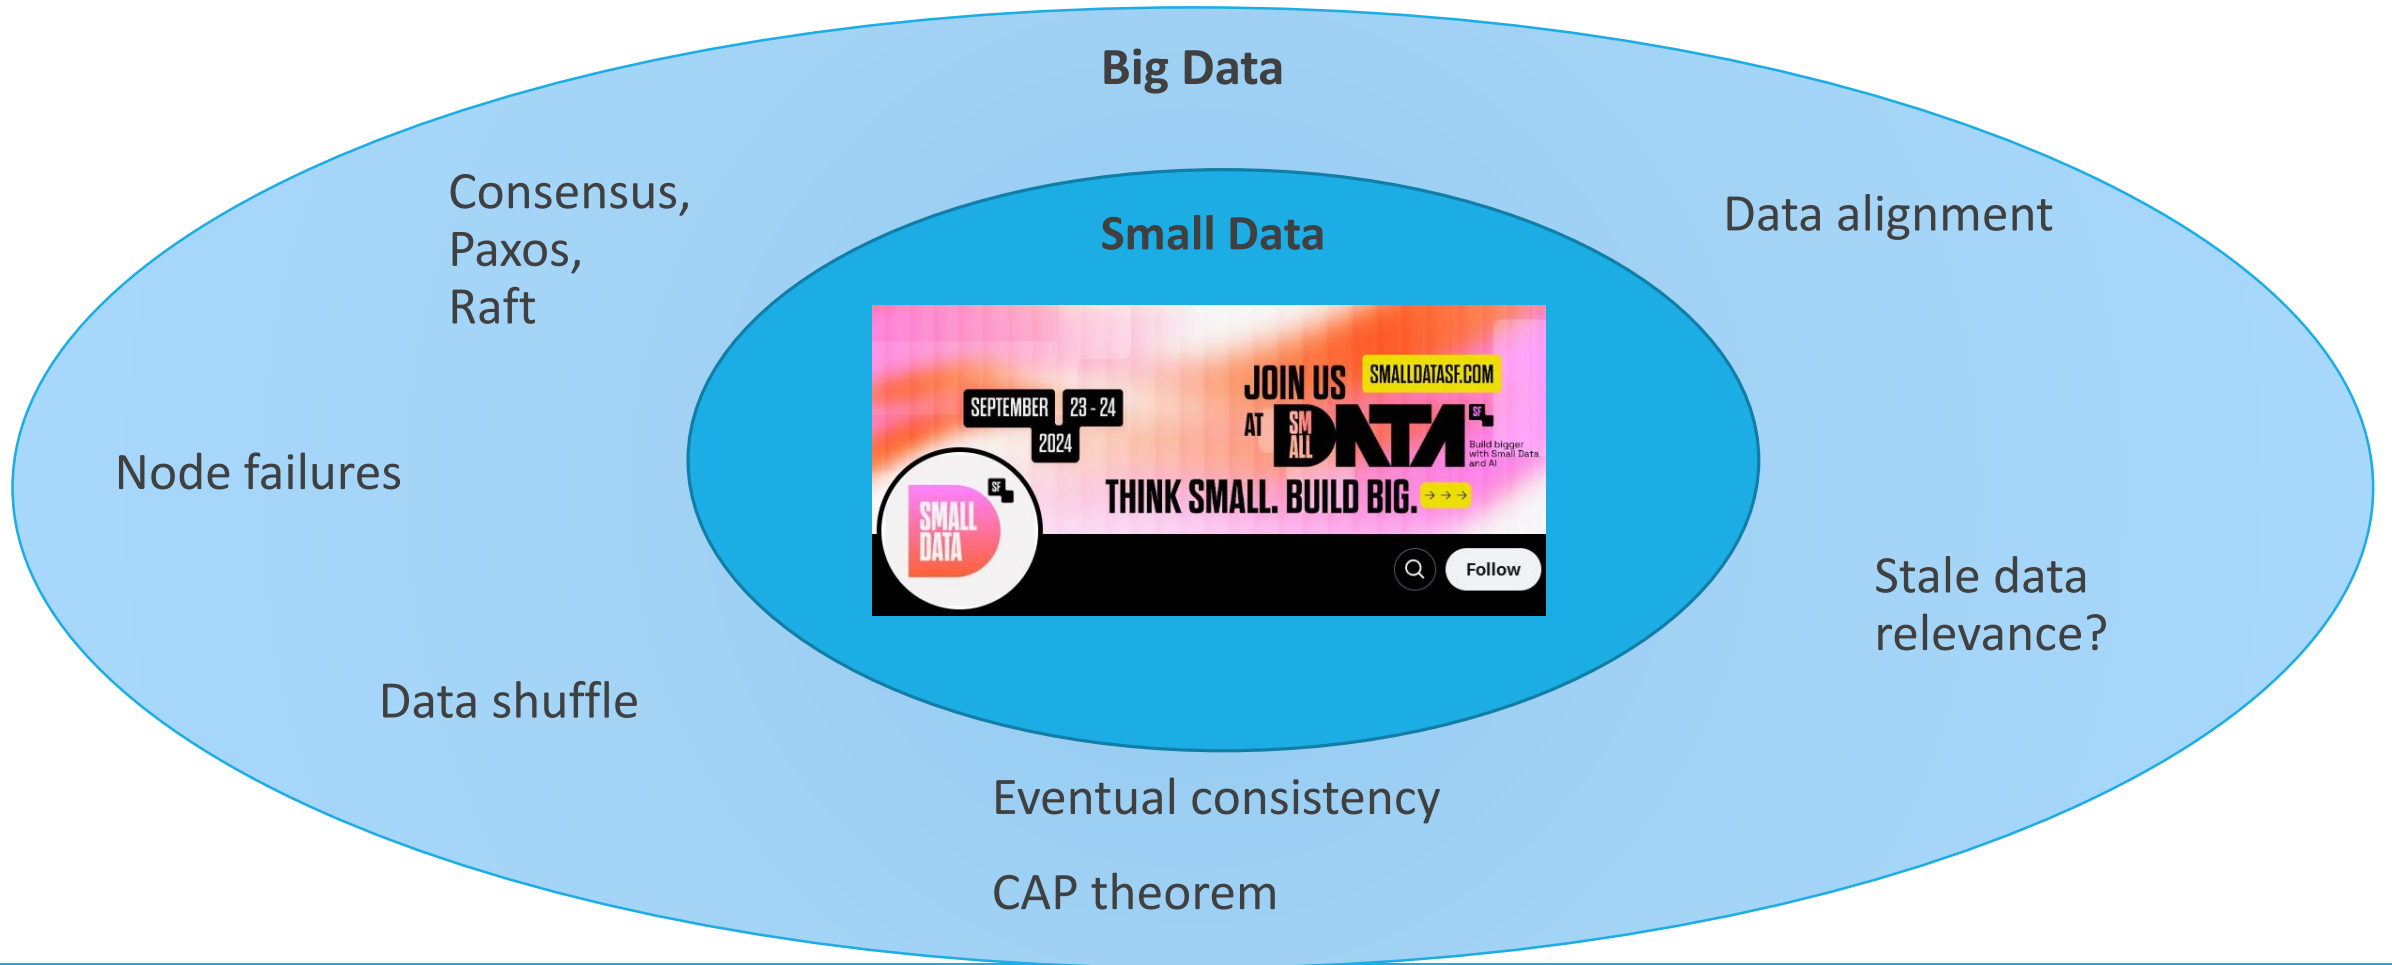
Exponential growth of hardware

Amazon EC2 offering

- 2006 - m1.small – 1 CPU + 1.7GB RAM
- 2024 - x8g.metal-48xl – 192vCPU + 3TB RAM
- 2024 - High-Memory (U-1) – 448vCPU + 24TB RAM



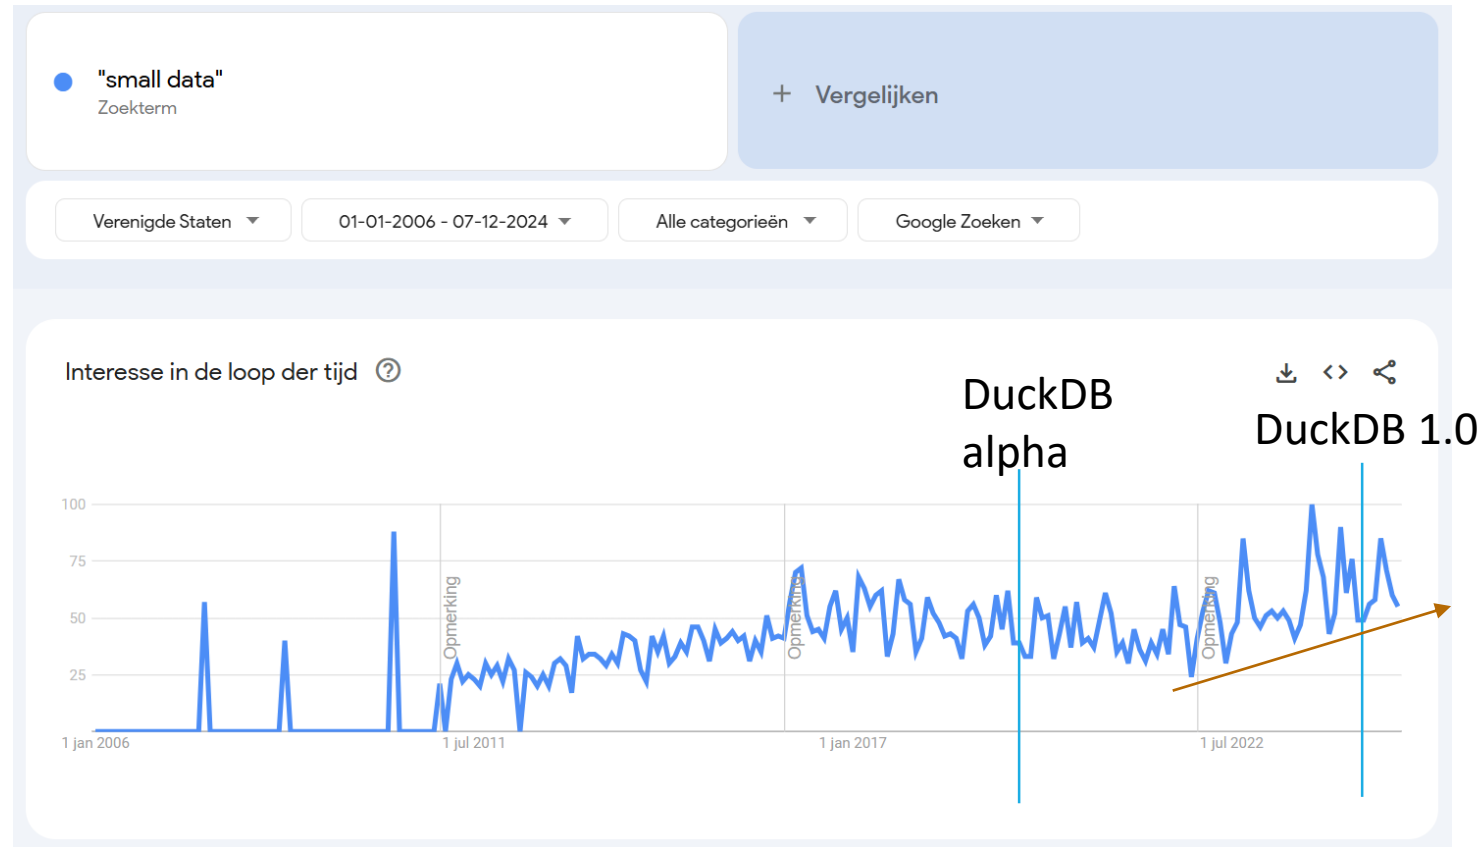
Rise of 'small data' movement



Google trends “big data”

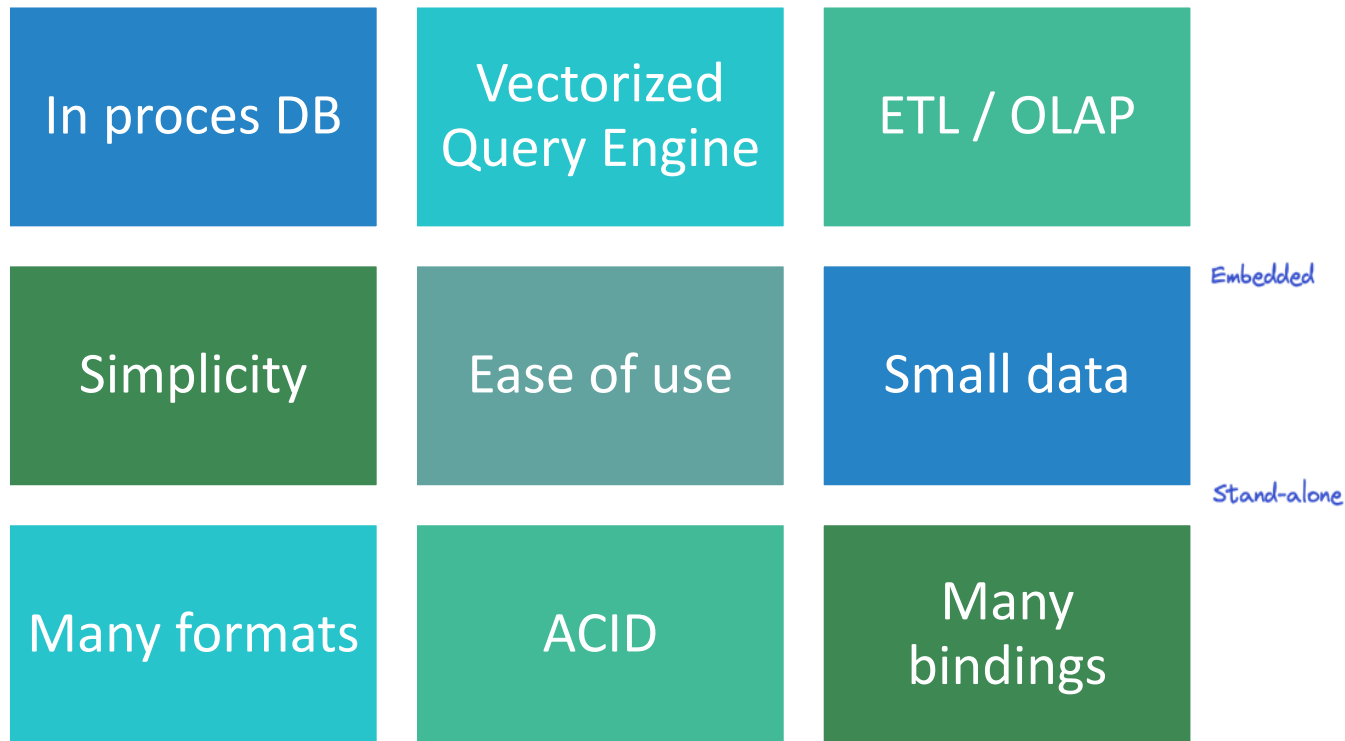


Google trends “small data”



The DuckDB sales pitch

"DuckDB: The SQLite for Analytics – Power, Simplicity, and Portability."



Transactional	Analytical
SQLite, SolidDB	DuckDB
Postgres, MySQL	Snowflake, ClickHouse, Redshift

Experiments

Research questions

OLTP vs OLAP using TPC-DS

- How does DuckDB perform against a traditional RDBMS under analytical workload conditions?

Scalability using TPC-H

- How does DuckDB scale with growing database size?

Comparative analysis using db-benchmark

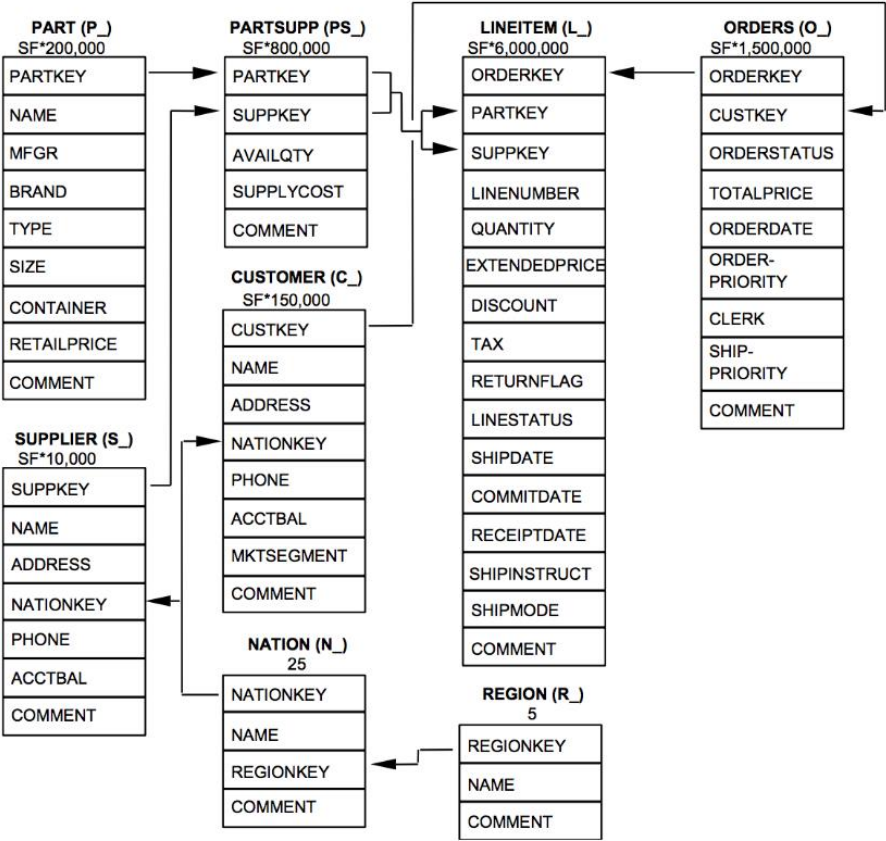
- How does DuckDB hold up against other in-memory OLAP processing systems?

source:

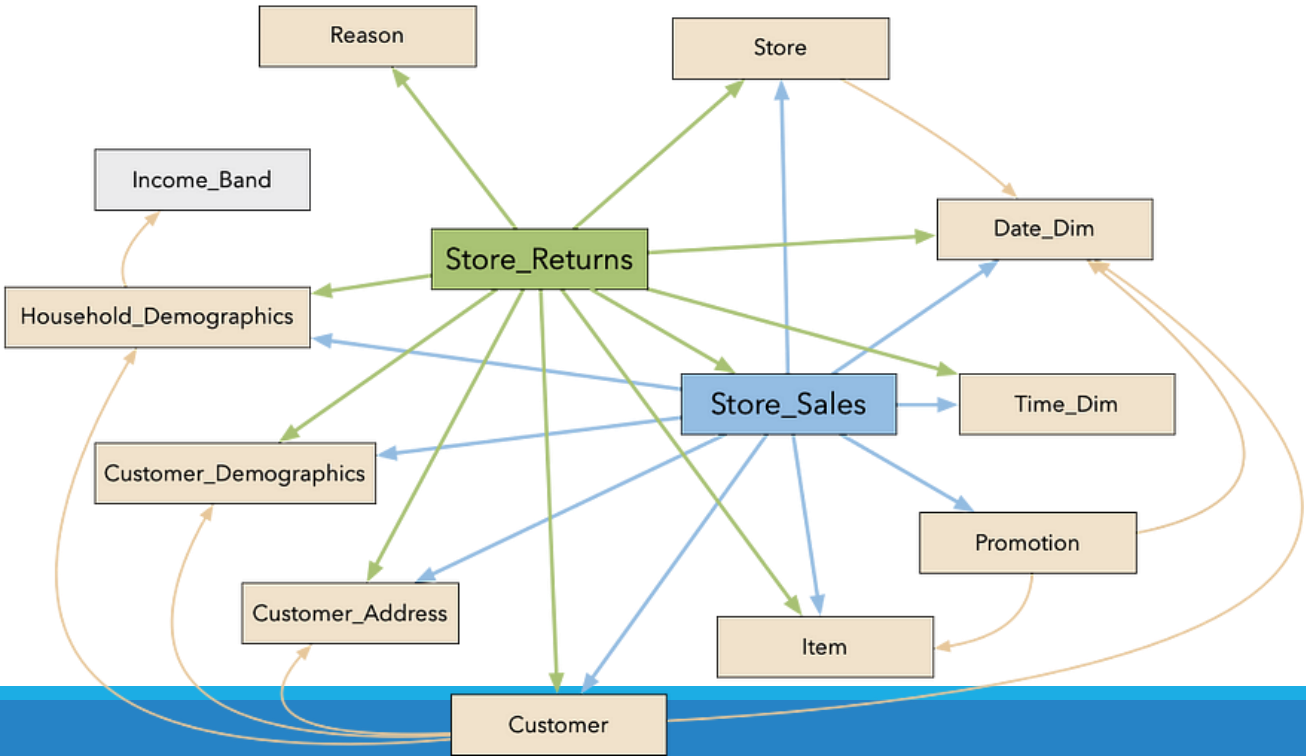
<https://docs.snowflake.com/en/user-guide/sample-data-tpch>

<https://medium.com/hyrise/a-summary-of-tpc-ds-9fb5e7339a35>

TPC-H vs TPC-DS



	TPC-H	TPC-DS
Schema type	3rd Normal Form	Multiple Snowflake
Number of tables	8	24
Number of columns (min)	3	3
Number of columns (max)	16	34
Number of columns (avg)	~ 7.6	18
Number of foreign keys	9	104



OLTP vs OLAP




Figure 3 Speedup of DuckDB over Postgres for the 84 TPC-DS queries that did not time out on postgres.

Query	Postgres [s]	DuckDB [s]	Ratio
1	timeout	0,097	
2	3,748	0,141	27
3	6,891	0,082	84
4	timeout	1,589	
5	4,945	0,201	25
6	timeout	0,162	
7	4,175	0,117	36
8	1,358	0,067	20
9	13,401	0,369	36
10	64,713	0,067	966
11	timeout	0,981	
12	0,536	0,042	13
13	3,056	0,122	25
14	timeout	1,356	
15	1,176	0,052	23
16	timeout	0,031	
17	timeout	0,099	
18	4,094	0,190	22
19	3,271	0,074	44
20	2,224	0,053	42
21	10,645	0,037	288
22	158,019	1,838	86
23	156,759	1,112	141
24	5,430	0,171	32
25	timeout	0,077	
26	4,080	0,096	42
27	4,674	0,250	19
28	28,235	0,295	96
29	11,100	0,142	78
30	timeout	0,123	
31	43,010	0,162	265
32	250,402	0,011	22764
33	6,754	0,074	91


Query	Postgres [s]	DuckDB [s]	Ratio
34	4,232	0,101	42
35	timeout	0,305	
36	8,791	0,250	35
37	18,750	0,043	436
38	45,684	0,262	174
39	108,574	0,280	388
40	3,121	0,056	56
41	33,772	0,031	1089
42	4,282	0,041	104
43	8,448	0,089	95
44	16,202	0,118	137
45	1,506	0,095	16
46	6,077	0,168	36
47	58,687	0,528	111
48	12,484	0,149	84
49	17,403	0,129	135
50	8,067	0,175	46
51	33,327	2,090	16
52	3,868	0,049	79
53	4,718	0,078	60
54	timeout	0,103	
55	3,911	0,048	81
56	8,690	0,111	78
57	24,801	0,359	69
58	9,338	0,090	104
59	37,343	0,289	129
60	8,755	0,127	69
61	7,974	0,125	64
62	6,321	0,063	100
63	4,446	0,069	64
64	18,913	0,545	35
65	26,490	0,415	64
66	5,969	0,183	33

Query	Postgres [s]	DuckDB [s]	Ratio
67	85,954	4,063	21
68	4,034	0,200	20
69	17,325	0,149	116
70	22,054	0,189	117
71	7,351	0,246	30
72	44,301	0,414	107
73	3,864	0,106	36
74	timeout	0,842	
75	28,367	0,446	64
76	9,813	0,162	61
77	9,212	0,093	99
78	26,832	1,856	14
79	4,845	0,180	27
80	11,426	0,328	35
81	timeout	0,262	
82	19,376	0,080	242
83	1,169	0,089	13
84	0,929	0,099	9
85	5,123	0,124	41
86	4,751	0,121	39
87	47,340	0,379	125
88	32,602	0,359	91
89	6,456	0,164	39
90	2,888	0,027	107
91	0,735	0,036	20
92	148,510	0,047	3160
93	7,959	0,275	29
94	timeout	0,089	
95	timeout	0,730	
96	4,541	0,035	130
97	14,309	0,381	38
98	5,216	0,494	11
99	14,206	0,099	143

Scalability

File name	Scale Factor (GB uncompressed)	Size on disk in MB (compressed)
tpch-sf1.db	1	254
tpch-sf3.db	3	771
tpch-sf10.db	10	2.612
tpch-sf30.db	30	7.965
tpch-sf100.db	100	26.962

Table 1 Overview of the TPC-H databases used for the benchmark.



```
SELECT
  c_count,
  count(*) AS custdist
FROM (
  SELECT
    c_custkey,
    count(o_orderkey)
  FROM
    customer
  LEFT OUTER JOIN orders ON c_custkey = o_custkey
  AND o_comment NOT LIKE '%special%requests%'
GROUP BY
  c_custkey) AS c_orders (c_custkey,
  c_count)
GROUP BY
  c_count
ORDER BY
  custdist DESC,
  c_count DESC;
```



Figure 3 Query execution time for first run (black) and second run (gray) on TPC-H with SF=100

Scalability (absolute and relative)

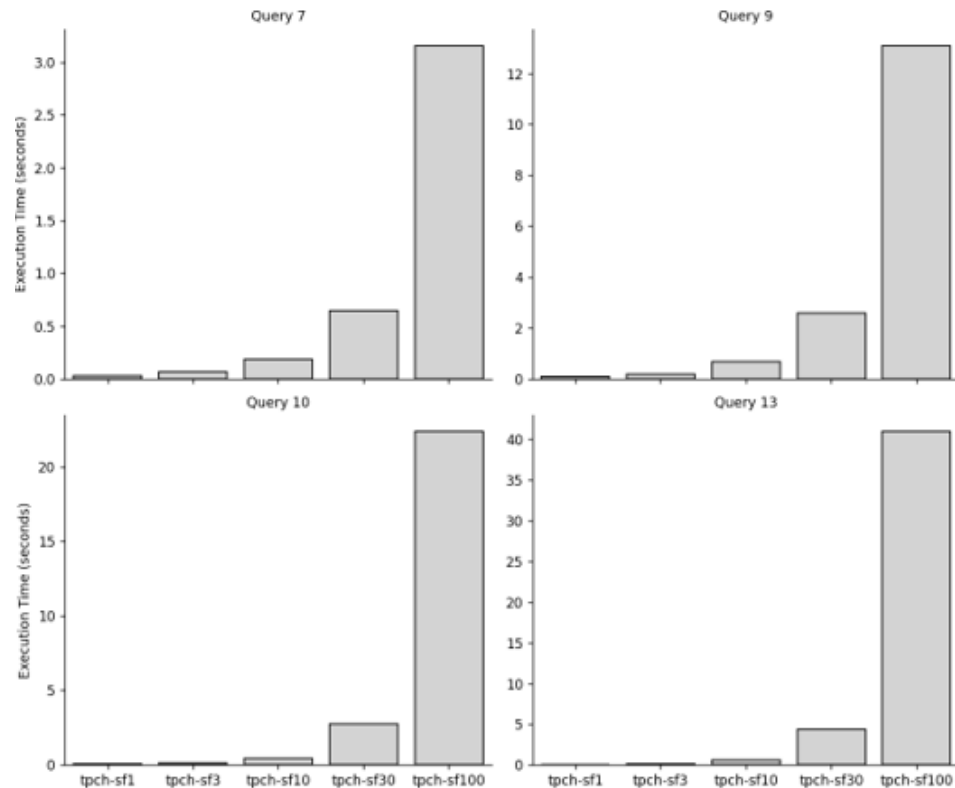


Figure 4 Query execution time for 4 selected queries on different TPC-H databases.




Figure 5 TPC-H query execution time divided by database size for 4 selected queries.

Comparative Analysis

<https://duckdblabs.github.io/db-benchmark/>

10 groupby queries

5 join queries

dataset size = 0,5 ; 5 ; 50 GB (.csv)

2 runs: focus on the first

runs on EC2 c6id.metal

- CPU: Intel(R) Xeon(R) Platinum 8375C CPU @ 2.90GHz
- CPU cores: 128
- RAM GB: 250

famous benchmark in DS/DA land

Task

groupby join

0.5 GB 5 GB 50 GB

basic questions

Input table: 1,000,000,000 rows x 9 columns (50 GB)

ClickHouse	24.8.4.13	2024-09-13	27s
DuckDB	1.1.0	2024-09-25	32s
Datafusion	41.0.0	2024-09-17	42s
Polars	1.8.2	2024-09-30	46s
data.table	1.16.99	2024-09-13	87s
collapse	2.0.16	2024-09-13	225s
spark	3.5.2	2024-09-13	243s
R-arrow	17.0.0.1	2024-09-13	403s
pandas	2.2.2	2024-09-13	822s
(py)datatable	1.2.0a0	2024-09-13	918s
dplyr	1.1.4	2024-09-13	1194s
InMemData.jl	0.7.21	2024-09-30	CSV import Segfault
DataFrames.jl	1.6.1	2024-09-30	CSV import Segfault
dask	2024.9.0	2024-09-17	out of memory
Modin		see README	pending


First time
Second time

Minutes 0.5 1.0 1.5 2.0 2.5 3.0 3.5

Query 1: "sum v1 by id1": 100 ad hoc groups of ~10,000,000 rows; result 100 x 2			
clickhouse	SELECT id1, sum(v1) AS v1 FROM tbl GROUP BY id1		
	0.00; 0.00		
duckdb	SELECT id1, sum(v1) AS v1 FROM tbl GROUP BY id1		
	0.01; 0.01		
R-arrow	AT %>% group_by(id1) %>% summarise(v1=sum(v1, na.rm=TRUE))		
	0.01; 0.01		
datafusion	SELECT id1, SUM(v1) AS v1 FROM x GROUP BY id1		
	0.01; 0.01		
polars	DF.groupby('id1').agg(pl.sum('v1')).collect()		
	0.01; 0.01		
spark	SELECT id1, sum(v1) AS v1 FROM tbl GROUP BY id1		
	0.07; 0.04		
collapse	collap(x, v1 ~ id1, sum)		
	0.08; 0.07		
data.table	DT[, .(v1=sum(v1, na.rm=TRUE)), by=id1]		
	0.11; 0.09		
dask	DF.groupby('id1', dropna=False, observed=True).agg({'v1': 'sum'}).compute()		
	0.13; 0.04		
pandas	DF.groupby('id1', as_index=False, sort=False, observed=True, dropna=False).agg({'v1': 'sum'})		
	0.29; 0.29		
dplyr	DF %>% group_by(id1) %>% summarise(v1=sum(v1, na.rm=TRUE))		
	0.30; 0.27		
pydatatable	DT[, {'v1': sum(f.v1)}, by=f.id1]		
	0.78; 0.62		
IMD.jl	combine(gatherby(x, :id1, stable = false), :v1 => IMD.sum => :v1)		
	CSV import Segfault: JuliaLang#55765		
DF.jl	combine(groupby(DF, :id1), :v1 => sum=skipmissing => :v1)		
	CSV import Segfault: JuliaLang#55765		

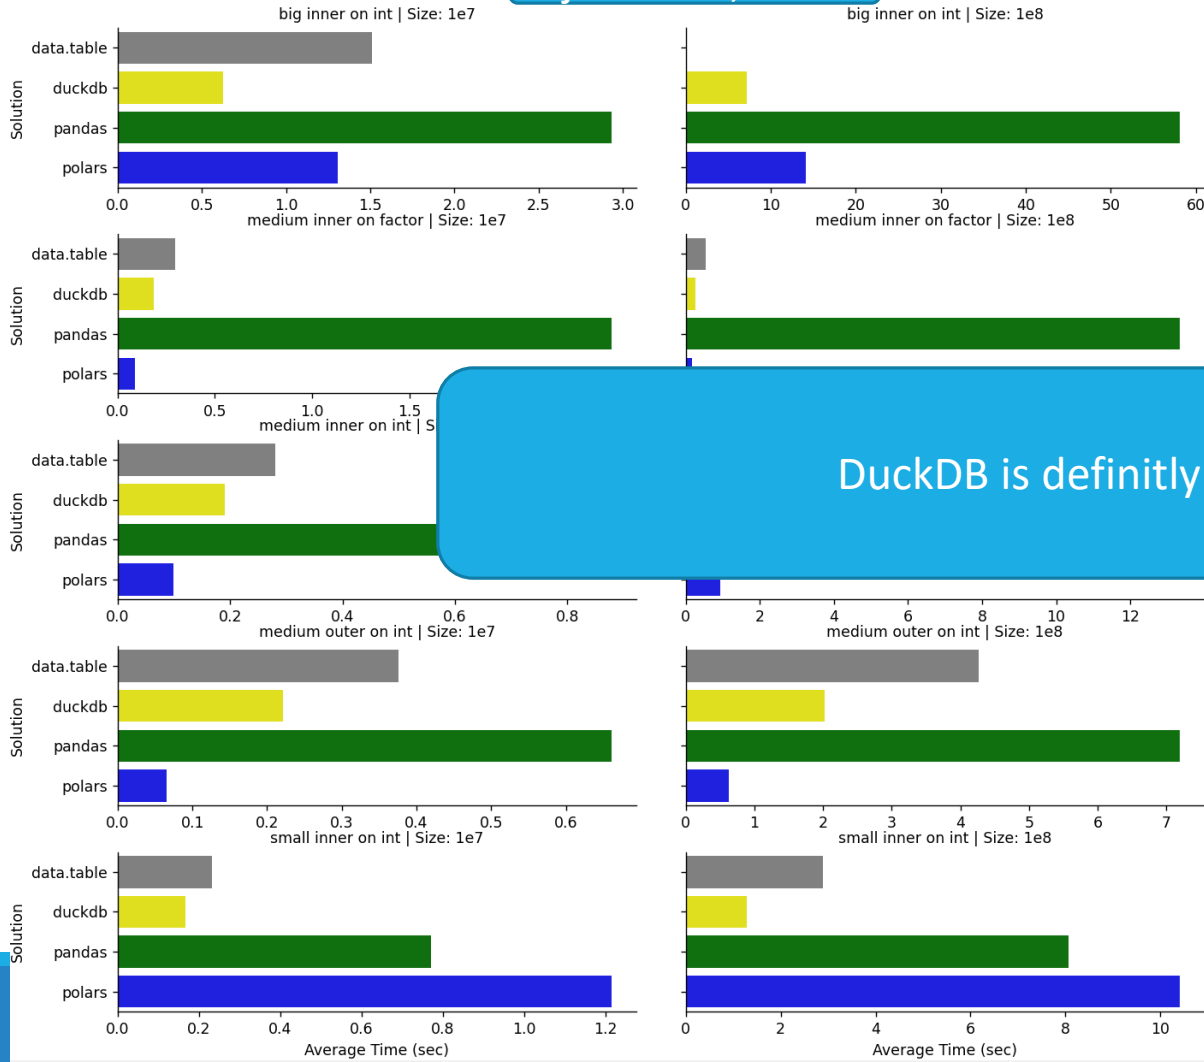
Conclusions

run	1				
in_rows	100000000				
Average of time_sec		Column Labels			
Row Labels		data,table	duckdb	pandas	polars
groupby		3,2875	1,6685	13,8044	5,8665
largest two v3 by id6		4,088	3,276	27,488	4,315
max v1 - min v2 by id3		4,036	2,295	4,802	2,018
mean v1:v3 by id4		1,65	0,088	1,483	8,644
median v3 sd v3 by id4 id5		4,362	0,946	4,881	3,618
regression v1 v2 by id2 id4		3,316	0,414	12,069	3,961
sum v1 by id1		0,952	0,044	1,039	3,439
sum v1 by id1:id2		1,177	0,188	3,218	4,498
sum v1 mean v3 by id3		1,849	2,66	7,279	9,106
sum v1:v3 by id6		1,487	2,583	7,402	7,66
sum v3 count by id1:id6		9,958	4,191	68,383	11,406
join		3,5845	2,7566	35,2052	5,4499
big inner on int			7,2065	58,076	14,1635
medium inner on factor		3,6635	1,6815	89,3655	1,1145
medium inner on int		3,519	1,582	13,326	0,941
medium outer on int		4,2625	2,02	7,1935	0,6285
small inner on int		2,893	1,293	8,065	10,402
Grand Total		3,4195	2,21255	24,5048	5,6582



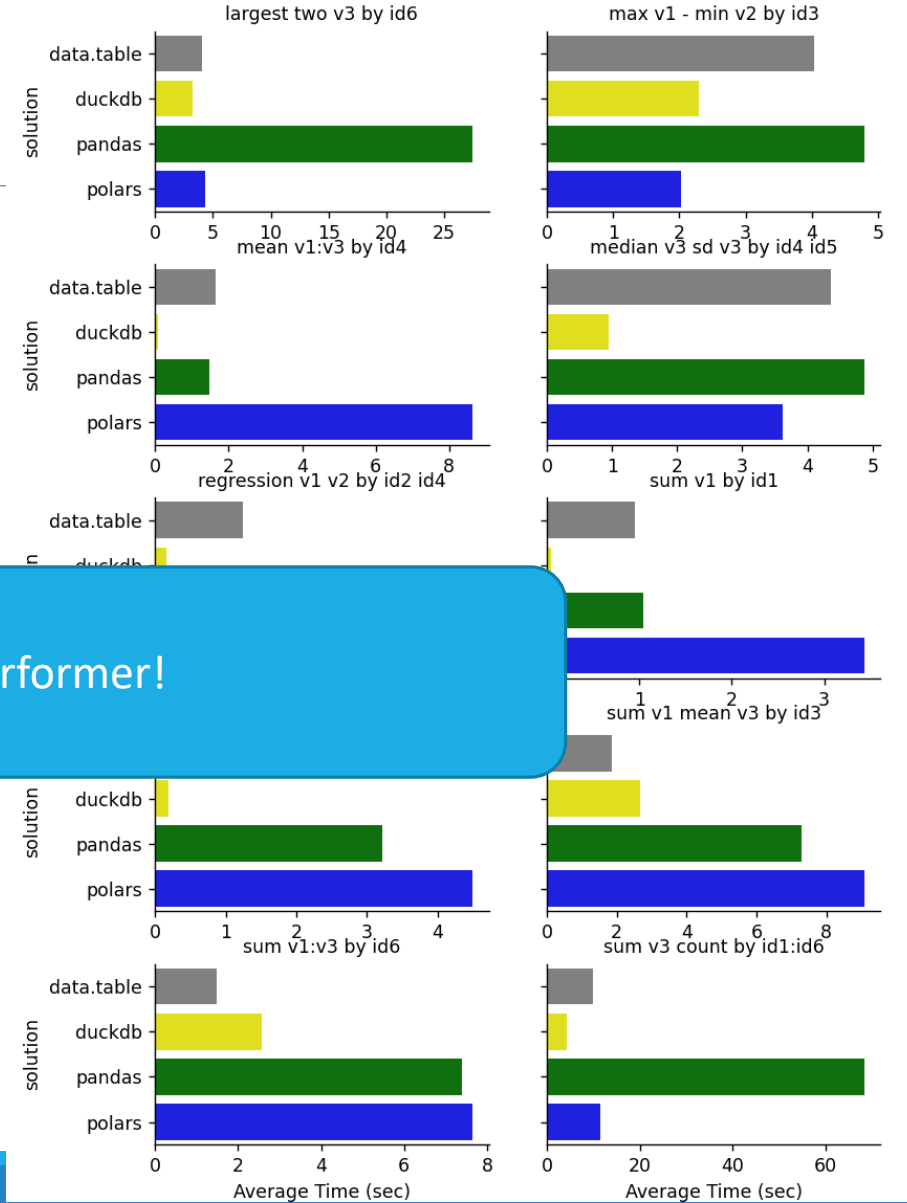
Conclusion

join 1e7 ; 1e8



DuckDB is definitely a top performer!


groupby 1e8



Conclusion

join 1e7 ; 1e8

groupby 1e8



PSA: DO NOT TRY
DB-BENCHMARK
AT HOME!



Take home messages

- ❑ “Small data” is still big – hardware is impressive – try to upscale before outscaling!
- ❑ If your DBA refuses your analytical queries, he has good reason too – use an OLAP query engine!
- ❑ Analysis on 100GB datasets can easily be done on single workstations
- ❑ Ease-of-use, available bindings, API stability, documentation, ... matter too. But it's always nice to be fast ;-)

Q&A
